Insert your title here

Insert a subtitle here (if applicable)



Presented by:
Mohohlo Samuel Tšoeu

Prepared for:
M. S. Tšoeu
Dept. of Electrical and Electronics Engineering
University of Cape Town

Submitted to the Department of Electrical Engineering at the University of Cape Town in partial fulfilment of the academic requirements for a Bachelor of Science degree in ***INSERT DEGREE NAME HERE (Electrical Engineering; Mechatronics or Electrical and Computer Engineering)

September 30, 2026

Declaration

1. I know that plagiarism is wrong. Plagiarism is to use another's work and pretend that it is one's own.
2. I have used the IEEE convention for citation and referencing. Each contribution to, and quotation in, this report from the work(s) of other people has been attributed, and has been cited and referenced.
3. This report is my own work.
4. I have not allowed, and will not allow, anyone to copy my work with the intention of passing it off as their own work or part thereof.

Signature:.....
M. S. Tšoeu

Date:.....

Terms of Reference

The terms of reference page is an agreement between yourself and your supervisor outlining what is expected of you in your final year project. Please make sure that this is discussed and written at the beginning of your thesis project.

Acknowledgments

Make relevant acknowledgements to people who have helped you complete or conduct this work, including sponsors or research funders.

Abstract

- Open the **Project Report Template.tex** file and carefully follow the comments (starting with %).
- Process the file with **pdflatex**, using other processors may need you to change some features such as graphics types.
- Note the files included in the **Project Report Template.tex** (with the .tex extension excluded). You can open these files separately and modify their contents or create new ones.
- Contact the latex namual for more features in your document such as equations, subfigures, footnotes, subscripts & superscripts, special characters etc.
- I recommend using the **kile** latex IDE, as it is simple to use.

Contents

1 Introduction **1**

1.1 Background to the study 1

1.2 Objectives of this study 1

1.2.1 Problems to be investigated 1

1.2.2 Purpose of the study 1

1.3 Scope and Limitations 2

1.4 Plan of development 2

2 Literature Review **3**

2.1 Wind Measurement in the Wind Energy Sector 3

2.1.1 Introduction 3

2.1.2 Current In-Situ Anemometry Methods and Limitations 4

2.2 LiDAR as an Alternative 4

2.2.1 Introduction 4

2.2.2 NML for Load Reduction 5

2.2.3 NML for Performance Optimization 5

2.2.4 LiDAR for Performance Assessment 6

2.2.5 Discussion of NML Applications 6

2.3 Principles of Doppler Wind LiDAR 7

2.3.1 Aerosol Backscatter and the Doppler Shift 7

2.3.2 Challenges in Measuring Doppler Shift in DWL 8

2.3.3 LiDAR Detection Modes 8

2.3.4 Coherent Detection: Homodyne, Heterodyne and Directional Ambiguity 11

2.3.5 Coherent Detection: Noise Sources and the Shot-Noise Limit . . . 13

2.3.6 Coherent Detection: Single-Ended vs Balanced Detection 15

2.3.7 Transmit and Receive Optics: Monostatic vs Bistatic 17

2.3.8 Ranging: Pulsed versus Continuous Wave 17

3 System Design **19**

3.1 Introduction 19

3.2 System Architecture 20

3.2.1 Coherent detection 20

3.2.2 Homodyne mixing 20

3.2.3 Balanced detection 21

3.2.4 Monostatic optics 21

3.2.5 Continuous-wave operation 22

3.2.6 All-fibre construction . . . 22
3.2.7 System Architecture Diagram . . . 22
3.3 System Requirements Analysis . . . 23
3.3.1 Functional requirements . . . 23
3.3.2 Performance requirements . . . 24
3.4 Component Choices . . . 24
3.4.1 Laser Source . . . 25
3.5 Final Design . . . 29

4 Noise Analysis **30**

4.1 Introduction . . . 30
4.2 A: Electronic, Shot and Relative Intensity Noise . . . 30
4.2.1 Experimental Goals . . . 30
4.2.2 Experimental Methodology . . . 31
4.2.3 Results and Discussion . . . 33
4.3 B: Phase-induced Intensity Noise . . . 38

5 Results **39**

5.1 Simulation Results . . . 39
5.2 Experimental Results . . . 39

6 Discussion **40**

7 Conclusions **41**

8 Recommendations **42**

A Appendix A **46**

A.1 Linewidth, PIIN and Doppler peak broadening 46

B Addenda **50**

B.1 Ethics Forms 50

Chapter 1

Introduction

1.1 Background to the study

A very brief background to your area of research. Start off with a general introduction to the area and then narrow it down to your focus area. Used to set the scene [**?**].

1.2 Objectives of this study

1.2.1 Problems to be investigated

Description of the main questions to be investigated in this study.

1.2.2 Purpose of the study

Give the significance of investigating these problems. It must be obvious why you are doing this study and why it is relevant.

1.3 Scope and Limitations

Scope indicates to the reader what has and has not been included in the study. Limitations tell the reader what factors influenced the study such as sample size, time etc. It is not a section for excuses as to why your project may or may not have worked.

1.4 Plan of development

Here you tell the reader how your report has been organised and what is included in each chapter.

Factors for architectural choices discussed in lit review Factors for component level design choices discussed in design

I recommend that you write this section last. You can then tailor it to your report.

Chapter 2

Literature Review

2.1 Wind Measurement in the Wind Energy Sector

2.1.1 Introduction

Wind has become a mainstream source of electricity. A record 165 GW of new wind capacity was installed worldwide in 2025, 40% more than the previous record year, bringing cumulative capacity to roughly 1.3 TW across 138 countries [9]. Delivering that growth at acceptable cost depends on how well the wind resource itself can be characterised, which Veers et al. identify as one of the grand challenges of wind energy science [22]. Wind speed is an important factor in analysis of a turbine's economics and engineering. The power available in the wind scales with the cube of wind speed, so a small error in measured or predicted wind speed becomes a much larger error in predicted energy [4]. Furthermore, wind speed and turbulence significantly affect the structural fatigue on a turbine, and act as an input to controllers which control the pitch and yaw of the turbines appropriately to maximise power generation and minimise stresses on the structure. It is clear that real-time measurement of the wind field around wind-turbines is a critical objective for this industry, and one to which much attention must be given. However, turbines operate in the lower atmospheric boundary layer, where the flow is turbulent and shaped by mesoscale circulation, orography, thermal stratification and the wakes of neighbouring turbines, which makes this characterisation of the wind-field difficult [17].

2.1.2 Current In-Situ Anemometry Methods and Limitations

The traditional instrument for wind-speed measurement for wind-turbines is the meteorological mast (met-mast), which holds cup and sonic anemometers at fixed heights, mounted on a long mast. This technology is well established, providing high resolution measurements, and indeed acting as a reference point to which any alternative method is compared [18], [12]. This methodology however holds several limitations:

Erecting a meteorological mast can be costly, particularly for offshore wind-turbines [24]. Once constructed, the met-mast is naturally limited to a single location, which limits spatial coverage - this makes variability in the wind, which occurs over larger areas, difficult to characterise [17]. Furthermore, the wake of the mast itself can corrupt measurements in the anemometer [17]. These constraints motivated the research into a remote sensing alternative to in-situ measurements.

2.2 LiDAR as an Alternative

2.2.1 Introduction

Advancements in LiDAR technology have resulted in the technology presenting an alternative and/or complement to existing wind-speed measurement techniques. This allows for the development of new turbine control and assessment strategies, which overcome some of the limitations of traditional wind-speed measurement techniques.

Doppler wind LiDAR (DWL) systems measure the line-of-sight velocity of wind by emitting a laser beam into the atmosphere and measuring the Doppler shift of the beam after being backscattered by naturally occurring aerosols and particles (see section 2.3). There are numerous commercial DWL systems available, such as Qinetiq's ZephIR whose first deployment was on the nacelle of a 2.3 MW turbine, and with Vaisala's (formerly Leosphere) pulsed WindCube shortly afterwards [12] [5]. Research has been focused on two areas; ground-based systems for wind-speed measurement (primarily for resource assessment prior to building wind-energy farms), and nacelle-mounted LiDAR (NML) systems [17].

Letizia et al. identified several promising application areas for NML systems, including load reduction, performance optimization, and performance assessment [17]. NML systems

present several advantages over traditional measurement techniques for these ends, including larger field of view and spatial coverage (particularly for scanning LiDARs) and reduced costs and maintenance [17]. The following subsections provide a brief overview of different NML application areas, and progress and limitations in these sectors. These areas of research present viable means to lower the levelized cost of energy (LCOE) for wind-turbines, and are thus well-suited to the growing demands for wind-energy systems globally [17].

2.2.2 NML for Load Reduction

Wind gusts can place significant load on a wind-turbine structure - however appropriate control of the turbine's rotor pitch can serve to decrease this load, preserving the structural integrity and longevity of the system. Standard turbine controllers are feedback systems that react to a turbulent gust only after it has loaded the rotor. A lidar preview of the approaching wind-field however would enable feed-forward pitch action **before** the gust arrives, and there is consensus that load reduction in certain regions of wind-turbine operation is a promising lidar-assisted control (LAC) application [17]. Crucially however, any benefit from feedforward LAC for load-reduction is limited by the coherence between the wind-speed measured by the LiDAR from a distance, and the actual wind-speed which hits the turbine rotors (rotor-effective wind-speed) [17]. The seminal simulation study of Harris et al. reported a 10% reduction in damage-equivalent load on the flapwise blade root [11], which spurred on further research into the topic. There is much ongoing research on the topic, particularly with respect to the evolution of the wind-field in between sensing and reaching the turbine, and reported results span a wide range. This lack of consolidation has thus presented a barrier to widespread adoption of LAC for load reduction, however research thus far indicates that this is indeed a promising application of the technology [21] [17].

2.2.3 NML for Performance Optimization

Another application area for LAC is performance optimization, where feedforward control based on NML measurements of the incoming wind-field can be used to yaw and pitch the turbine appropriately, and use torque control to maintain optimal rotor speed, so as to maximise power generation. This has however showed limited results in comparison to load reduction. Simulation studies thus far have found negligible gains in power coefficient from feed-forward torque control for optimising rotor speed, because conventional non-

LAC torque controllers are already mature [17]. Where gains are reported they come from yaw correction - a hub-mounted ZephIR used to correct misalignment on a 2 MW turbine gave an estimated 5% annual energy production increase [17].

2.2.4 LiDAR for Performance Assessment

Power performance testing is an area in which NML has had greatest success and widespread adoption [17]. A turbine's power curve is conventionally measured against a met-mast's wind-speed measurements. Mikkelsen et al. used a spinner lidar (mounted on the wind turbine blades rather than the nacelle) to reconstruct the low-speed end of a 2.3 MW turbine's power curve and found less scatter than the met-mast measurement, which they attributed to a more representative rotor-equivalent wind speed [19]. Wagner et al. achieved similar results, and identified the reasons: because the lidar yaws with the turbine it always measures upwind, so its measured wind-speed correlates better with power fluctuations than the met-mast's; and because it is never in the wake of the turbine under test, the usable wind sector was larger than the undisturbed sector of the cup anemometer [23]. This shows a direct advantage of a DWL system over traditional in-situ measurements, where the met-mast is fixed. A subsequent uncertainty analysis found the NML estimated annual energy production only 1.1% more uncertain than the cup-anemometer reference, its higher instrumental uncertainty offset by greater data availability [24]. The International Electrotechnical Commission (IEC) formalised the method in 2022 [13], and some 500 campaigns per year are projected this decade [17].

2.2.5 Discussion of NML Applications

The three application areas discussed above present results which show some of the use-cases of LiDAR for wind-speed measurement, however they are not exhaustive, and the topic remains an active area of research. What is however clear in the literature is that the widespread adoption of such systems depends on the reliability, accuracy and cost of the LiDAR systems [17]. Traditionally, such systems have been very expensive, and there has been research into methods to reduce the cost [3]. The remainder of this report thus addresses aspects related to the design of such systems, and the costs and tradeoffs involved.

2.3 Principles of Doppler Wind LiDAR

Doppler wind LiDAR uses a simple working principle in order to determine wind-speed from a distance. The following section details the physics and mathematical concepts which are used in DWL systems.

2.3.1 Aerosol Backscatter and the Doppler Shift

LiDAR systems emit very high-frequency electromagnetic radiation (when compared to its technological cousin - radar), thus, wavelengths of the emitted radiation will be very small. This property makes the technology uniquely suitable for the purposes of wind-speed measurement; the wavelength of the emitted LiDAR beam is comparable to the size of the aerosol particles, and thus the particles scatter some of the emitted beam (Mie scattering) [19].

The backscattered light can be detected by the LiDAR system's photodetector. The backscattered light, however, will be shifted in frequency relative to the incident LiDAR beam, by a magnitude which depends on the velocity of the particle itself - this is due to the Doppler effect [8]. Specifically, the shift in frequency of the backscattered light is directly proportional to the line-of-sight (LOS) velocity of the particle - that is, the component of the particle's velocity in the direction of the line of the LiDAR beam - through the following equation [17]:

$$v_{LOS} = \frac{c}{2}\frac{f_{beam} - f_{scatter}}{f_{scatter}} \tag{2.1}$$

Note the factor of 2 in the denominator, which is necessary to account for the fact that the Doppler shift occurs twice - once for the emitted light hitting the moving particle, and again for the light being emitted/scattered **from** the moving particle.

Thus, by determining the Doppler frequency shift of the received backscatter from aerosol particles, the LOS velocity of those particles can be determined, and from this, the LOS wind-speed can be inferred. Note that this explanation simplifies concepts used in ranging of the wind-speed, which will be discussed in later sections.

2.3.2 Challenges in Measuring Doppler Shift in DWL

A crucial challenge in DWL is the measurement of the Doppler shift itself. LiDAR beam frequencies are of the order of 10^{14} Hz (hundreds of terahertz) while, from equation 2.1, the Doppler shift in this THz-range signal, produced by aerosols moving at typical wind speeds is of the order of only 10^6–10^7 Hz (a few megahertz) - this is a fractional frequency shift of roughly one part in ten million which needs to be resolved. Photodetectors can not follow typical LiDAR optical signals, which are in the order of THz. For example, for a LiDAR system with beam wavelength 1550 nm, the field oscillates at 193 THz, orders of magnitude beyond the bandwidth of photodetectors or typical ADCs [1, 12]. Thus, the return backscatter signal's frequency can **not** be directly measured, and its frequency relative to the emitted beam simply observed from this. Instead, specialised techniques need to be used to recover the Doppler shift of the signal, **without** having to measure the frequency of the THz-range return signal for comparison. This is thus a central difficulty in DWL systems, and the detection architectures used to address it are discussed in the sections that follow.

2.3.3 LiDAR Detection Modes

LiDAR systems are typically grouped into two main types, each with their own distinct advantages, based on how received light is detected - these types are Direct Detection, and Coherent Detection. To varying extents, both of these detection modes have been applied in DWL systems, with significant differences in how the Doppler shift is measured in each. This section gives an overview of the principles of each detection mode, and discusses how this can be utilised to measure the Doppler shift of received backscatter light.

Direct Detection

Direct detection LiDAR systems operate on a very simple principle: A laser pulse is fired from the system; this pulse then gets reflected off a target (or backscattered in the case of aerosols), and the reflection is incident **directly** on the LiDAR system's photodetector. Owing to its simplicity, direct detection LiDAR systems are widely used in robotics and automotive applications for ranging and basic velocity measurements. The time between the emitted pulse and received pulse can be used to calculate the distance of targets from the system, and repeated distance measurements can be used to estimate the velocity of the target. Obtaining velocity through Doppler shift measurements however presents a

challenge to this architecture.

As discussed in Section 2.3.2, photodetectors by themselves can not be directly used to determine the Doppler shift, as they can not follow THz frequency light signals. Multiple modified direct detection DWL system architectures have been developed to address this problem. Most common is the "double-edge technique" illustrated by Korb et al. [16]. Two Fabry-Pérot Interferometers (etalons) tuned symmetrically around the transmit laser frequency are used to split the backscatter signals - and thereafter the Doppler shift may be obtained from the ratio between the intensity of the two signals, as measured by photodetectors [7].

Crucially for this method, the ratio between scattering types (Rayleigh and Mie) for the aerosols, as well as temperature must be known in order to produce accurate measurements - this can vary significantly with altitude [25]. These factors, and strict requirements on etalon tuning, temperature, and laser stability present a considerable challenge in the implementation of this architecture [7]. In the double-edge technique any drift of the laser relative to the etalon edge is indistinguishable from a Doppler shift, so the transmit laser must be absolutely stable to roughly 100KHz of its nominal frequency, which in practice requires an active lock to an atomic transition [7]. The etalon itself must be temperature controlled to about 1 milli-Kelvin, since a 1 milli-Kelvin drift produces a $0.95m/s$ error, and mechanical drift of the mount can produce biases of $1.5 - 2m/s$ per week, requiring weekly recalibration [7].

An advantage that direct detection architectures hold over coherent detection is the ability to use molecular backscattering return (Rayleigh scattering), which gives it better performance in high-altitude regions with less aerosols to give particulate backscattering (Mie scattering, which a coherent detection system relies on) [7]. However, direct detection DWL systems typically have worse accuracy and spatial resolution than their coherent detection counterparts, which are more technologically mature [25]. For these reasons, Direct detection architectures are not as common in the literature on DWL, as compared to coherent detection systems, particularly in typical wind-turbine applications in the aerosol-rich boundary layer, where the primary advantage of direct detection architectures is unnecessary. The remainder of this review will thus focus on coherent detection systems, and the intricacies of their system architecture.

Coherent Detection

Coherent detection LiDAR varies from direct detection in that the return signal (from backscatter or reflections) is optically mixed with a local oscillator (LO) signal prior to detection. Typically, a single laser source is used, where a set fraction of the beam is emitted at the target, and the other fraction is used as the LO for mixing down the line. Coherent detection provides multiple distinct advantages over direct detection, which will be described in this section.

Firstly, the optical mixing of the signals makes it possible to electronically distinguish a MHz-range Doppler shift on top of a THz-range carrier (as was identified to be a key challenge in DWL in Section 2.3.2). This can be described mathematically as follows: If the laser source (and therefore LO) has frequency f_0, and the return signal has frequency $(f_0 + \Delta f)$, where Δf is the frequency shift of the return signal due to the Doppler effect, the signal after mixing the return signal with LO is given by:

$$E_{mixed}(t) = \sqrt{2P_{LO}}\cos(2\pi f_0 t) + \sqrt{2P_{RS}}\cos(2\pi(f_0 + \Delta f)t) \tag{2.2}$$

Where P_{LO} and P_{RS} are the time-averaged optical powers of the local oscillator and return signal respectively.

The photodetector current, I_{det}, is proportional to the intensity of the electromagnetic field of the light signal - $|E_{mixed}|^2$. When expanded, one obtains:

$$\begin{aligned} I_{det}(t) \propto\ & \underbrace{P_{LO} + P_{RS}}_{\text{DC}} + \underbrace{2\sqrt{P_{LO}P_{RS}}\cos(2\pi\Delta f\, t)}_{\text{beat}} \\ & + \underbrace{P_{LO}\cos(4\pi f_0 t) + P_{RS}\cos(4\pi(f_0 + \Delta f)t) + 2\sqrt{P_{LO}P_{RS}}\cos(2\pi(2f_0 + \Delta f)t)}_{\text{not followed by detector}} \end{aligned} \tag{2.3}$$

It can be seen from the above equation that the squared mixed signal has three components: one at DC, one with frequency in the order of f_0, and another with frequency at Δf. Note again that f_0 is in the THz range, where Δf is in the MHz range. As discussed, if a MHz-range bandwidth photodetector is used to detect E_{mixed}, the detector cannot follow oscillations at frequencies in the order of f_0, so those terms produce no response

and may be omitted. Thus, Equation 2.3 simplifies to the following equation, which is called the beat signal:

$$I_{det}(t) \propto [P_{LO} + P_{RS}] + 2\sqrt{P_{LO}P_{RS}}\cos(2\pi\Delta f\, t) \tag{2.4}$$

Thus, if the frequency of this beat signal is measured, the Doppler shift, Δf, is obtained. In this way it can be seen that coherent detection allows for direct measurement of the Doppler shift, and therefore the velocity of the target.

Aside from the direct measurement of the Doppler shift, the "signal beating" process described above provides several other key advantages. The beating process amplifies the return beam, which allows for shot-noise limited operation (the ideal case for such a system), and extremely high sensitivies [12] (elaboration in 2.3.6). This is essential in DWL, as return signals in clear atmospheric conditions are typically very low, due to low backscatter from aerosols. Furthermore, this detection process rejects signals which are not common in spatial mode and polarisation with the LO, meaning that sunlight, or other LiDAR beams will not affect the measurement [12].

2.3.4 Coherent Detection: Homodyne, Heterodyne and Directional Ambiguity

The following section discusses two distinct classes of coherent detection methodology, and how the choice of class affects the recovery of the wind direction from the signal.

Homodyne

While not explicitly stated, the previously discussed direct down-conversion of the return signal to the beat signal - via mixing with a local oscillator with the **same** frequency as the emitted signal (Equation 2.4 and preceding) - is known as **homodyne** coherent detection. Homodyne coherent detection down-converts the MHz-range doppler shift in the THz range carrier to a MHz-range beat signal around baseband, i.e. the frequency of the signal is directly the frequency of the Doppler shift. This has significant implications:

Ignoring the DC components in Equation 2.4, we get

$$I_{det}(t) \propto \cos(2\pi\,\Delta f\,t)\,. \tag{2.5}$$

The Doppler shift (Δf) therefore appears directly at baseband. However, because the cosine is an even function,

$$\cos(-2\pi\,\Delta f\,t) = \cos(2\pi\,\Delta f\,t)\,, \tag{2.6}$$

Therefore, noting that Δf is signed based on whether the target is approaching or receding, it can be seen that approaching and receding targets of equal speed produce identical detector currents. Only the magnitude $|\Delta f_D|$ is recoverable, and the measured spectrum is symmetric about zero frequency. Note that Abari et. al[1] demonstrated a methodology for obtaining signed velocity, even with a homodyne coherent detection setup, however this setup is novel and introduces additional complexities and costs which are outside of the scope of this project.

Heterodyne

The alternative to this is known as **heterodyne** coherent detection. In heterodyne architectures, a slight frequency shift is imposed on the laser - crucially however this shift is only applied to one of the fractional outputs, either the local oscillator **or** the emitted laser (but not both). This means that when the same beating process described previously is applied, the MHz-range doppler shift in the THz range carrier is down-converted to a MHz-range beat signal centered around an intermediate frequency (f_i below), rather than baseband, as with homodyne architectures. This is conventionally achieved with an acousto-optic modulator, placed either in the local oscillator arm [10] or in the transmit arm [2]. This has the advantage of allowing for the recovery of the target's direction:

Taking the offset in the local oscillator arm, such that $f_{LO} = f_0 - f_i$, the beat term becomes

$$I_{det}(t) \propto \cos(2\pi\,(f_i + \Delta f)\,t)\,. \tag{2.7}$$

Provided $f_i > |\Delta f|_{\max}$, the sign of Δf is recoverable. Approaching targets appear above f_i and receding targets below it, and the sign of the velocity is recovered.

2.3.5 Coherent Detection: Noise Sources and the Shot-Noise Limit

In a coherent detection DWL system, there are four primary sources of noise which need to be considered, namely electronic noise, shot noise, relative-intensity noise, and phase-induced intensity noise:

Electronic noise arises from the photodetector and the transimpedance amplifier (which converts photodetector current into a measurable voltage), and includes dark noise (signal present when there is no light incident on the detector) and thermal noise [12]. This noise is constant, independent of LO power [5], [2].

Shot noise is a fundamental noise source in coherent detection. Shot noise arises from the discrete nature of photons, which reach the photodetector at slightly different random instants [20], i.e. it is a property of light itself, and can therefore not be removed. Because the photodetector's current is generated by these received photons, the detector then produces a white, Gaussian noise current whose power spectral density is proportional to the intensity of the light incident on it [20]. In the case of a coherent detection setup, it is thus proportional to P_{LO}. The significance of shot noise follows from comparing its scaling with that of the beat signal. From Equation 2.4 the beat current has amplitude $\propto 2\sqrt{P_{LO}P_{RS}}$, so the signal power is proportional to P_{LO}. Shot noise power is also proportional to P_{LO} [6, 2]. Electronic noise power, however, as mentioned, is constant with P_{LO}. Neglecting other noise sources for the moment, the SNR is approximately [6]:

$$\text{SNR} = \frac{P_{LO}P_{RS}}{P_{EN} + P_{SN}} = \frac{P_{LO}P_{RS}}{P_{EN} + \alpha P_{LO}} \tag{2.8}$$

Where P_{EN} is electronic noise power, P_{SN} is the shot noise power, and α is a proportionality constant representing the other terms which multiply with P_{LO} to give the shot-noise power. Observing thus that P_{SN} grows proportionally to P_{LO}, it can be seen that at low LO powers the electronic noise term dominates, and the SNR will grow as P_{LO} scales (because the numerator scales with P_{LO} as well). At very large LO powers, the electronic noise power term becomes negligible relative to the shot noise, and thus the SNR simplifies to:

$$\text{SNR} \approx \frac{P_{LO}P_{RS}}{P_{SN}} = \frac{P_{LO}P_{RS}}{\alpha P_{LO}} = \frac{P_{RS}}{\alpha} \tag{2.9}$$

This is called ideal **shot-noise limited** operation, and it is the best possible operating point for a coherent detection system - the LO has raised the signal (and shot noise) so far above the electronic noise floor that the electronic noise floor is swamped out so as to become negligible, and no further sensitivity is available from the receiver. The SNR reaches a constant peak which is proportional only to the received backscatter signal power (P_{RS}), and the constant α which governs shot-noise power. [3]. In practice however, the ideal shot-limited operating point can not be fully reached, because two effects prevent P_{LO} from simply being made as large as possible: the photodetector saturates [2], and the laser's intensity noise, discussed next, grows faster than the signal.

Relative-intensity noise (RIN) is fluctuation of the source laser's relative intensity, which appears as current noise at the output of the detector [12], creating a wideband noise floor in the frequency domain. As can be derived from Equation 2.4, the useful detectable signal **power** grows with P_{LO}. However, RIN power grows with P_{LO}^2, and thus grows faster than signal power [2]. Recall from above that we desire a large P_{LO} in order to achieve shot-noise limited operation, with larger P_{LO} giving a better SNR in Equation 2.8 **when RIN is not included**, and the ideal case being an infinitely large P_{LO}. With RIN **included** however, as in:

$$\text{SNR} = \frac{P_{LO}P_{RS}}{P_{EN} + P_{SN} + P_{RIN}} \tag{2.10}$$

It can be seen that beyond a certain P_{LO} level RIN will overtake shot noise in the denominator, and the system's SNR will no longer be increasing with increasing P_{LO} and approaching the ideal shot-limited regime [2]. Instead, the SNR will **decrease** with an increasing P_{LO} after this crossover point, due to the P_{LO}^2 scaling of RIN [6]. It can thus be seen that RIN establishes an optimal P_{LO} which can be used, and sets a ceiling on how close the detector can get to the ideal shot-noise limited operation of Equation 2.9, before increasing RIN dominates. It is therefore clear that if RIN can be reduced through other means, a larger P_{LO} can be used, and thus the system can get closer to the ideal shot-noise limited state. Balanced detection is one means to achieve this [6], and is discussed in subsection 2.3.6.

Finally, there is phase-induced intensity noise (PIIN), which is noise which results from non-ideal linewidth effects of the laser source [6], [14], [10]. PIIN is a significant source of noise in DWL systems, and has been shown to possibly exceed the shot-noise by up to 30dB [10]. Thus, in order to remain as close as possible to the ideal shot-noise limit, PIIN should be reduced. This is possible through the selection of an appropriate laser source,

and through careful matching of signal path lengths in the DWL system [10][6]. PIIN is a complex noise source which can vary significantly in shape and magnitude depending on the laser source and the DWL system. An understanding of laser phase noise and linewidth is required to understand it - a full discussion and mathematical breakdown on these concepts and how they relate to PIIN is provided in section A.1.

2.3.6 Coherent Detection: Single-Ended vs Balanced Detection

As previously discussed, in order to determine the Doppler shift of the received backscatter signal in a coherent detection DWL system, the backscatter signal must be optically mixed with the LO signal, and this mixed signal must then be detected using a photodetector system. This section discusses the two common types of photodetection methodologies - single-ended detection, and balanced detection.

As discussed in subsection 2.3.5, the best-case performance of a DWL system is for it to operate in a shot-noise limited regime - thus, all other sources of noise should be reduced to below the shot-noise level. In coherent detection LiDAR architectures, the use of a balanced detector setup has become common for this purpose, as it assists in reducing a laser source's relative-intensity noise (RIN) below the shot-noise level, in a manner which single-ended detection does not.

A single-ended detection system simply uses a 2:1 fibre optic coupler to mix the LO reference signal with the backscatter signal, and uses a single photodetector to convert this light signal into a measurable current, which will be proportional to the optical signal power as given in Equation 2.4. It can be seen that this detection method retains the DC terms P_{LO}, P_{RS}. Any RIN from the laser source is present in the output, as these intensity fluctuations are contained in these two power terms. Thus, the measurable Doppler peak will be smaller with respect to the noise floor.

Contrastingly, balanced detection setup makes use of a 2:2 fibre optic coupler which also optically mixes the LO reference signal with the received backscatter signal (as discussed in section 2.3.3). However, the 2:2 coupler has two outputs, thus giving two output signals which can be termed E_1 and E_2 [15]:

$$E_1 = \frac{1}{\sqrt{2}}(E_{RS} + E_{LO}), E_2 = \frac{1}{\sqrt{2}}(E_{RS} - E_{LO}) \tag{2.11}$$

Where E_{RS} and E_{LO} are the complex electromagnetic fields of the returned backscatter

signal and the local oscillator respectively.

Thereafter, the two signals enter the balanced detector, which consists of two matched photodiodes and a transimpedance amplifier to convert their currents into a measurable voltage. The photodiodes respond with currents I_1 and I_2 respectively, proportional to the intensity of the electromagnetic field incident on it ($|E_1|^2$ and $|E_2|^2$ respectively), as was also discussed in section 2.3.3. Due to the internal setup of the transimpedance amplifier and the photodiodes in the balanced detector, the output current is then given by [15]:

$$I_{det} \propto I_1 - I_2 \propto |E_1|^2 - |E_2|^2 \propto 2\sqrt{P_{LO}P_{RS}}cos(2\pi\Delta ft) \tag{2.12}$$

Comparing this to Equation 2.4, it can be seen that the DC terms P_{LO} and P_{RS} are no longer present due to them being common-mode in both I_1 and I_2 and thus being canceled by the subtraction. This therefore lowers the noise floor, as the RIN contained in these terms is also eliminated.

Figure 2.1 depicts this process graphically:

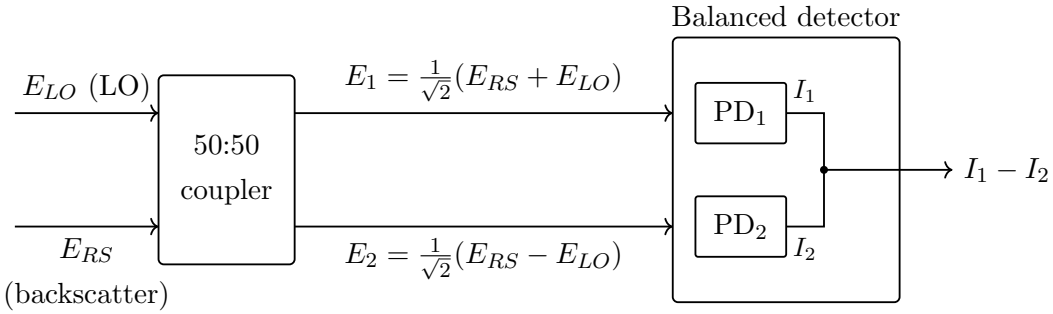


Figure 2.1: Balanced detection. The 50:50 coupler delivers the sum and difference of the backscatter and LO fields to two matched photodiodes; the output is the difference of their photocurrents, in which the common-mode terms cancel and only the beat term $2\sqrt{P_{LO}P_{RS}}\cos(2\pi\Delta ft)$ remains.

It can thus be seen that balanced detection provides the advantage of cancelling RIN. This allows for the final detected signal to theoretically be able to achieve the desired shot-noise limited operation (if other sources of noise are also eliminated).

2.3.7 Transmit and Receive Optics: Monostatic vs Bistatic

For a DWL system there are two optical configurations which are commonly used - bistatic and monostatic. This section discusses the differences between the two.

Bistatic

Bistatic configurations use one lens to emit the laser into the atmosphere, and another to receive the backscatter signal. One distinct advantage of a bistatic configuration, is that back reflections of the beam from the lens used for emission are not coupled into the mixed signal, and thus do not appear at the detector [6]. This serves to reduce phase induced intensity noise (discussed in section 3.4.1. However, this is at the cost of greatly increased complexity, due to additional considerations of alignment of the two lens apertures [10] [3].

Monostatic

A monostatic configuration uses a single lens aperture for both the emission of the laser, and the receiving of the backscatter signal. A fiber optic circulator can then be used to ensure the emitted signal and the received signal do not interfere, and are routed appropriately in the system. This has the distinct advantage of not requiring any additional lenses, and alignment considerations, and indeed has become the norm for such systems due to this reduced complexity [12].

2.3.8 Ranging: Pulsed versus Continuous Wave

Coherent DWL systems have two distinct methodologies which can be applied in order to determine at what distance from the system a measured wind-speed is occurring. The methodology differs depending on a key architectural choice of the system, namely whether the LiDAR beam is on for longer periods of time - known as continuous wave (CW) - or if it is pulsed. Both of these architectures hold their own advantages and disadvantages, which will be discussed in this section.

Pulsed LiDAR

As the name suggests, a pulsed LiDAR system emits its beam in a short pulse. Depending on the pulse duration, a given volume of air - called the probe volume - will be illuminated at an instant of time, and will create backscatter. This volume naturally moves through the air as the beam propagates radially. Thus, by using extremely precise timing on the returned backscatter, it can be determined which distance "slice" of air that backscatter was generated by, and thus, what the wind-speed for that particular "slice" is. The wind-speed for that distance is averaged over the probe volume, thus a shorter pulse will give greater spatial accuracy due to its shorter probe volume. This technique is known as range-gating, and has been widely applied in commercial and scientific DWL systems. [5] [23] [17]. A distinct advantage of range-gating is that the probe volume remains the same at all distances, including longer ones, giving it constant range resolution for longer distances than CW systems [5] (described below). However, it must also be noted that typical pulsed systems are ineffective in ranges below tens of metres, because the pulse length is at least as long as the desired range [12].

Continuous-Wave LiDAR

Continuous-wave systems emit a continuous beam through the atmosphere. This simplifies the design cost and complexity of the systems [5] [12], and no precise pulse timing or receive timing is required. CW systems achieve ranging by using optics to focus the beam at a specific distance. This gives a Lorentzian spatial weighting function along the beam length, with a peak at the focus point of the beam (the point with highest intensity), thus concentrating the detection sensitivity around the focus, at the expense of contributions from other ranges. However, the effective probe volume varies **quartically** with the distance of the focal point from the source, thus significantly increasing the volume over which the Doppler shift (and therefore wind-speed) is averaged at longer focal distances. Therefore for longer ranges, spatial resolution severely decreases, and it is not possible to gate the measurements for any specific region of air before the focal point. For this reason, CW systems are not well-suited for long ranges of several hundred meters [12], however at short ranges they can be more effective than pulsed systems, owing to the previously mentioned limitations of pulsed systems, relating to pulse length.

Chapter 3

System Design

3.1 Introduction

This chapter details the design of a Doppler wind LiDAR system. As discussed in the introduction and scope of this report, the system is a benchtop prototype whose purpose is to make the sensitivity of a DWL measurable and attributable to specific component and architecture choices. As such, the primary considerations will be noise-levels, sensitivity, and the overarching principles of physics governing DWL systems. This system design will thus not include the design of housing for field deployment, nor will it be concerned with real-time post-processing pipelines for the system's data (which would be required in a field-deployed system). Furthermore, several key components that would be typical in a scientific- or industry-grade, field-deployable DWL system are deliberately omitted in order to reduce costs, at little functional expense for a benchtop system. Where possible, the system was made extensible for future iterations and field deployment.

First subsection 3.2.1 gives an overview of the architecture chosen, as synthesised from the various architecture types reviewed in chapter 2, and with respect to the project's overarching goals. Thereafter, **??** discusses the functional and performance requirements of the system. Section 3.4 details the specific components chosen to meet these requirements. Finally, **??** presents two numerical simulations which validate the choice of components and architecture for the project.

3.2 System Architecture

From the reviewed literature in chapter 2, it was decided that an **all-fibre, coherent detection, continuous-wave, monostatic** DWL system with **balanced detection** would be built, with sampling and digitisation to be handled by a benchtop oscilloscope, and post-processing done on a laptop computer. This section details why these architectural choices were made, using summarised information from chapter 2 with respect to the project goals, and discusses any significant implications for the system.

3.2.1 Coherent detection

The direct detection technique requires the Rayleigh/Mie scattering ratio and temperature, an absolutely stabilised laser, etalon control to 1 milli-Kelvin, and very precise physical alignment and calibration [16, 7, 25]. These present substantial challenges to an undergraduate project, particularly the stabilised laser and etalon temperature control. The one advantage which direct detection holds, which is the capability of using molecular backscatter when aerosol backscatter is not available, is irrelevant in an aerosol-rich controlled lab environment, which is the only environment this system will be tested in.

By contrast, coherent detection is far more common in the literature, and thus has more practical information and sources for components. Furthermore, mixing with the LO also permits shot-noise-limited operation through the effective amplification of the return signal, and rejects unwanted light not mode-matched to the LO [12] (e.g. other LiDAR systems, or the sun), thus reducing sources of potential interference. The architecture thus lends itself to a more elegant, direct method of measuring the Doppler shift, via optical mixing and subsequent detection.

For these reasons, a coherent detection system was chosen for this project build.

3.2.2 Homodyne mixing

Homodyne mixing places the Doppler shift at baseband with no additional frequency shift, and no additional components necessary, whereas heterodyne operation requires an acousto-optic modulator (AOM) in the LO or transmit arm [10, 2]. The tradeoff is the sign ambiguity as discussed in subsection 2.3.4, where only the velocity **magnitude** is

measured in a homodyne system, whereas a heterodyne system can resolve the direction.

For a benchtop-scale system which primarily aims to present the working principles of a DWL system and a sensitivity analysis, and where the direction of the flow being measured will always be known, it is not necessary to recover signed velocity. Adding an AOM thus adds needless complexity, and adds an additional unnecessary cost to a budget-constrained project. Crucially however, the addition of an AOM for a future iteration of the system is simple, and does not require the redesign of any parts of the current system. For these reasons it was thus decided to omit the AOM, and to therefore construct a **homodyne** coherent detection system.

3.2.3 Balanced detection

As discussed in subsection 2.3.6, signal power scales with P_{LO} but RIN with P_{LO}^2 [2]. However, to approach shot-noise limited operation with respect to electronic noise, a large P_{LO} is desired, as outlined in subsection 2.3.5. Thus, at the LO power needed for ideal shot-noise-limited operation a single-ended photodetector will be RIN-limited at large P_{LO} values. A 2:2 coupler and matched photodiode pair (balanced detection) cancel the common-mode P_{LO} and P_{RS} terms in Equation 2.4 and with them the RIN [15]. Thus, in order to approach the ideal shot-noise limited operation, which helps increase sensitivity in a DWL system with weak backscatter returns [12], a balanced detector setup will be beneficial, and was thus chosen for use in this project.

3.2.4 Monostatic optics

For this project, it was decided to use a monostatic optical configuration - a single aperture serves both transmit and receive, separated by a fibre circulator, with no second lens and no overlap alignment necessary. A bistatic arrangement would keep transmit-optics back-reflections away from the detector, at the cost of greatly increased alignment complexity [6, 10, 3], as well as the additional cost of 2 sets of optics, and equipment needed for alignment. The accepted tradeoff here is thus that back-reflections from the transmit path reach the detector and raise the noise floor through PIIN (see section 3.4.1).

3.2.5 Continuous-wave operation

Pulsed systems are ineffective below tens of metres because the pulse length is typically at least as long as this range [12]. a CW system needs no pulse or receive timing, which simplifies design and can reduce costs (due to lower component count). It is also the more effective of the two at short range. One of the significant performance costs of a CW system however is reduced spatial resolution, particularly at long ranges, as spatial resolution is defined by focus rather than gating with timing (see **??**). Seeing as a benchtop system requires operation at short ranges in the order of tens of centimeters to one or two meters, a CW system will be best suited here due to its simplicity, and its advantages at short range, as well as the irrelevance of long-range limitations.

3.2.6 All-fibre construction

For this project it was decided that an all-fibre construction would be used - meaning that fibre optic cables would be used to transmit the LO and emission beams, and return signals, rather than transmitting the beams through free-space. Single-mode fibre makes the mode-matching between LO and return that coherent detection requires automatic, removes free-space alignment complexities, and gives access to widely and cheaply available telecom components [3, 1, 2]. Non polarisation-maintaining fibre is used, owing to cost considerations as compared to polarisation-maintaining fibre. The accepted tradeoff of standard (non-polarisation-maintaining) fibre is that the polarisation of LO and return is not fixed, which will reduce the heterodyne efficiency of the system, and thus reduce the measureable Doppler peak [6].

3.2.7 System Architecture Diagram

Figure 3.1 below shows a functional block-diagram of the system's decided-upon architecture. In this diagram, the signal path can clearly be seen from laser source, through the splitter (splitting into LO and transmit beam). Thereafter the transmit beam gets routed via port 1 –>2 of the circulator to the monostatic optics, which also receives the backscatter signal to be routed via port 2 –>3 of the circulator. The local oscillator and backscatter signals are mixed in the optical mixer, and then detected by a balanced photodetector, from which the final output beat signal can be sampled and processed.

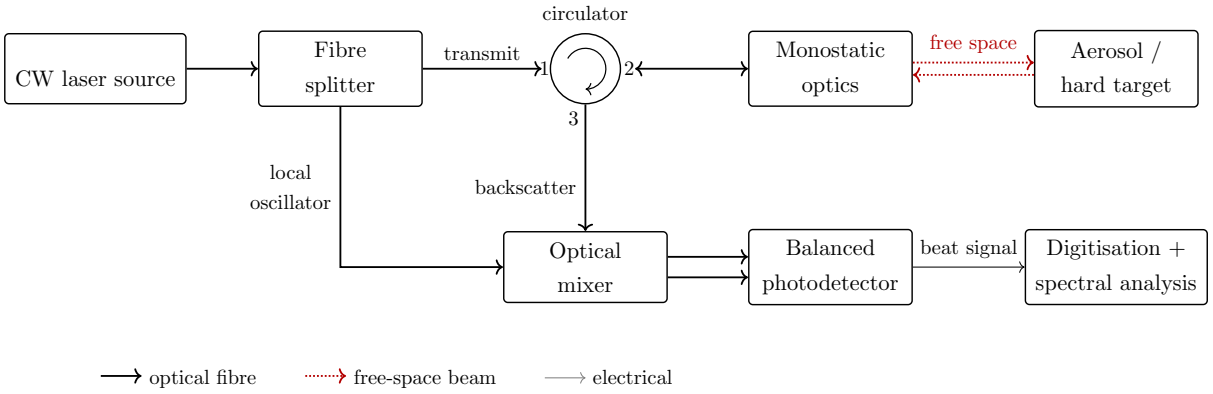


Figure 3.1: Block diagram of the chosen architecture: all-fibre, coherent, continuous-wave, monostatic, homodyne, with balanced detection

3.3 System Requirements Analysis

Now that a system architecture has been decided on, requirements which are specific to this choice of architecture, and which are related to the project goals can be synthesised. The requirements for the system are divided into two sets: functional requirements, which define the ideal outcomes which are desired from the system, and performance requirements, which define specific metrics which are to be measured and quantified with respect to the goals of the functional requirements.

3.3.1 Functional requirements

1. Measure the line-of-sight velocity of a target from the Doppler shift of coherently detected backscatter, using a monostatic all-fibre homodyne architecture with balanced detection and continuous-wave operation.

2. Operate on both a diffuse hard target of known velocity (for calibration) and on air seeded with smoke or fog.

3. Produce a velocity estimate from a recorded spectrum using a stated, reproducible detection criterion (threshold relative to the system's noise floor).

3.3.2 Performance requirements

From the desired functional requirements, a set of performance requirements can be synthesised; these are given in Table 3.1 below along with the IDs of the tests which will verify them. It must be noted again here that this project's goals chiefly concern the measurement and characterisation of noise and SNR in a constrained DWL system build, and how these aspects relate to component choice. As such, the majority of these requirements are concerned with the characterisation of the system relative to an ideal system, and thus the quantification of the system's limitations - rather than attempting to achieve performance specifications which will allow all the functional requirements to be met.

ID	**Requirement**	**Value**	**Verified by**
P1	Velocity detection on a hard target	Measured Doppler peak at least 10 dB above noise floor for a target 10cm away from laser emission	C1
P2	Detector noise	approaching ideal shot-noise-limited behaviour: total detector noise (electronic noise + residual RIN after balanced detection) within 1 dB of the shot-noise limit	A1
P3	Total system noise	total noise floor (shot + electronic + residual RIN + PIIN) quantified with respect to the shot-noise limit, and sources of excess noise described and characterised	B1
P4	Phase-induced intensity noise (PIIN)	level (in dB) and spectral shape measured; agreement with the path-mismatch model of [10] within 6 dB	B1, B2
P5	Low-velocity limit	minimum resolvable LOS speed quantified; goal of $\leq 0.1\,\mathrm{m\,s^{-1}}$ on a hard target	B1, C1
P6	Detection of wind-speed in seeded air (smoke or fog)	Doppler peak detected at $5 \times \sigma$ (5 times the noise standard deviation) above the mean, at $N \leq 1000$ averages **OR** measurements indicating noise-floor characteristics which make this not possible	D1, D2

Table 3.1: Performance requirements and their verification tests.

3.4 Component Choices

This section details the specific component selections which were made for the system. These selections were made to best suit the functional and performance requirements of

the system. It must be noted that this project is constrained when compared to typical scientific or industrial DWL systems. As such, many component choices were driven by cost, availability, and time-frame considerations, rather than simply their suitability for the system. This is however a suitable approach for the goals of the project, which are to illustrate the principles and physics involved in a DWL system, and to quantify system performance with respect to component choices.

3.4.1 Laser Source

The selection of an appropriate laser source is critical to ensure good performance of a DWL system. There are numerous properties of a laser source which should be considered, and which will be detailed in this section. After discussion on these properties, a suitable laser source is selected.

Wavelength

There are many suitable wavelengths for DWL systems, as detailed by Cariou et. al [6]. There is much discussion around SNR optimisation and backscatter maximisation for varying wavelengths [3], however owing largely to practical advantages, a wavelength of 1550nm has emerged as a common choice for both commercial and scientific DWL systems [1]. Fiber optics telecommunications commonly operate at this wavelength - thus making components designed for this wavelength readily commercially available and significantly cheaper than alternatives [6], [2], [12]. Additionally, laser beams at 1550nm are eye-safe at significantly higher power levels (up to multiple Watts), compared to shorter wavelengths [6], reducing safety concerns. For this reason it was decided that a 1550nm laser source would be used for this project.

Power Level

As previously discussed, the ability of a DWL system to measure wind-speed depends on the measurement of aerosol backscatter of the emitted beam.

The time-averaged power of the backscatter signal into a receiver can be expressed as [12]:

$$P_s = \pi P_t \beta(\pi) \lambda \tag{3.1}$$

Where P_s is the time averaged backscatter signal power, $\beta(\pi)$ is the atmospheric scattering coefficient (evaluated at π, i.e. scattering at π radians, back towards the source), and λ is the wavelength of the transmitted light. Note that $\beta(\pi)$ also holds a wavelength dependence, and thus P_s does not necessarily scale linearly with increasing wavelength, as the equation may suggest, but instead resolves to a near cancellation of the wavelength dependence between some frequency ranges [3].

It can be seen in Equation 3.1 that there is a linear dependence on P_t, the power of the transmitted beam. Thus, a higher-powered transmit laser source would give a larger return signal to be measured, increasing the sensitivity of the system by increasing the SNR of the system [3].

A setup commonly used to achieve higher emitted laser power here is the master-oscillator power-amplifier (MOPA) configuration [6], [3]. This separates the functions of the laser source - allowing for a low-power, highly coherent, narrow linewidth laser source to be used as a master-oscillator and an amplifier to increase the power levels of the master oscillator beam, while maintaining its spectral purity. This also allows the local oscillator to be derived from the master oscillator, eliminating the need to have two separate laser sources - one for LO, one for output - which would need to have a stable frequency relationship [6], something which requires careful design to achieve, as illustrated by Colberg et al. [7].

As noted by Bennett, such power amplifiers (commonly Erbium-doped fiber amplifiers; EDFAs) make up a significant portion of the cost of a DWL system. Indeed, more than 20 years after the writing of Bennett's paper, the cost of EDFAs remain prohibitively expensive for a project at this scale. For EDFA systems which are able to amplify signals to hundreds of milli-Watts or multiple Watts (which is a typical power level for such a system [2][5][6]), current prices range in the thousands to tens of thousands of USD (see Appendix).

For this reason, the use of a power amplifier for the laser source is disqualified for this project on the grounds of cost. The system was however designed such that the addition of such an amplifier would be possible at a later stage.

Due to the relatively low powers available for laser sources without power amplification (in the tens of milli-Watts), other avenues to improve detection sensitivity will be required.

Bennett identifies another option for increasing the sensitivity of a DWL system - which is to average over many returns [3]. This is the approach which will be taken for this project. Additionally, owing to this constraint in laser source power, all tests which concern wind-speed measurement in air will be performed using air seeded with smoke and/or fog. This will increase backscatter levels by increasing $\beta(\pi)$ in Equation 3.1, and thus reduces the need for the power amplifier for this benchtop-scale system.

Linewidth

Linewidth can be seen as a measure of the spectral purity of a laser source, and is one of the most critical parameters in the design of a DWL system. A narrow linewidth is highly desirable for DWL applications, as this will significantly reduce the phase-induced intensity noise (PIIN) in a DWL system, and allow the system to therefore have a better SNR and sensitivity to detect small backscatter signals. Furthermore, the linewidth of a laser source has a significant effect on how broad the returned backscatter signal's Doppler shifted peak is, particularly at longer ranges. This will result in the system having worse precision, and not being able to resolve turbulence in the air as well. For a full discussion on the mathematical and physical meaning of linewidth, as well as a treatment on how linewidth affects PIIN and the broadened peak, see section A.1.

Laser source types and selection

There are many types of laser sources available on the market. In the selection of a laser source for this project, the primary driver was cost and availability, with the design goals being to maximise power level and minimise linewidth, while fixing the wavelength at 1550nm. After selecting the best available laser source for these parameters, the performance of the laser source will then be analysed, and recommendations based on its properties and performance can then be made.

Source type	Linewidth (typ.)	Power, unamplified	Noise and practical notes	Indicative cost (USD)
Fabry–Pérot laser diode	multi-mode; not single-frequency	10 mW to 100 mW	Multiple longitudinal modes; signal from different modes interferes destructively beyond the coherence length [10]. Unsuitable for coherent detection.	< 100
DFB laser diode (butterfly, TEC)	0.5 MHz to 5 MHz; ∼100 kHz in newer devices [6, 10]	10 mW to 40 mW	Relaxation peak ∼1 GHz, far above the Doppler band; RIN flat and low in-band [6]. Coherence length tens of metres in fibre. Needs external driver and TEC controller.	100 to 1000 (module)
Planar external-cavity semiconductor (e.g. RIO ORION/PLANEX)	$<$ 1 to 15 kHz	10 mW to 30 mW	Gain chip with planar Bragg-grating cavity; RIN shot-noise-limited above 500 kHz; integrated driver; marketed for wind lidar.	quote only; low thousands
External-cavity diode laser (grating, Littrow/Littman)	10 kHz to 100 kHz	10 mW to 50 mW	Narrowed by optical feedback, but prone to jitter and mechanical instability [10]. Free-space cavity.	quote only; ∼10 000+
Er-doped DFB fibre laser (e.g. NKT Koheras BASIK)	$<$ 0.1 to 10 kHz [10]	10 mW to 40 mW	The source of choice for CW anemometry at range [6, 12]; 5.2 kHz unit in [2]. Relaxation peak near 1 MHz, inside the low-velocity Doppler band, so balanced detection is required there [6].	quote only; tens of thousands
MOPA (any seed + EDFA)	that of the seed	0.1 W to 10 W	Required for pulsed and long-range CW operation [6, 1]; adds the amplifier's noise figure and ASE.	seed + 3600 to 13 000 (benchtop EDFA); quote for $>$1 W

Table 3.2: Single-frequency laser source classes at 1550 nm. Linewidths and powers are typical of commercial devices; costs are list-price tiers as of September 2026 (see **??**) and exclude drivers, controllers and import costs.

3.5 Final Design

The final design of the system using the architecture and components discussed above is given in Figure 3.2 below.

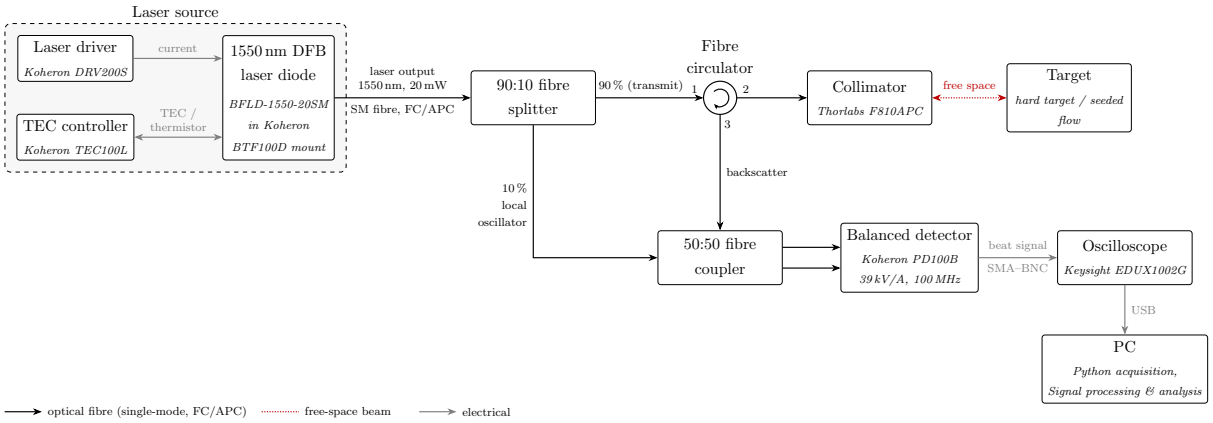


Figure 3.2: Architecture diagram of final system design

Chapter 4

Noise Analysis

4.1 Introduction

This chapter gives a comprehensive account of the testing which was done on the completed DWL system to account for the various sources of noise. The purpose of the noise analysis is to quantify the noise observed in the system and attribute it to the appropriate sources, while investigating how component, architecture, and operating point choice affects this. This chapter is separate from the functional testing of **??**, as the experimental setup differs significantly, and the results from these experiments will indeed be used to inform the expected results of the experiments in **??**, and to construct an optimal experimental setup in those cases.

4.2 A: Electronic, Shot and Relative Intensity Noise

4.2.1 Experimental Goals

As discussed in subsection 2.3.5, a coherent DWL performs best when it is shot-noise limited: the shot noise of the local oscillator (LO) must dominate every other noise source at the detector. Experiment A characterises the receiver's noise floor with the LO as the only light on the detector, so that the individual noise contributions can be separated and measured. It has the following goals:

- **Measure the electronic noise floor (A1).** With the laser off, measure the noise of the PD100B and acquisition chain alone, and compare it with the input-referred noise specified in the PD100B datasheet. Agreement with the datasheet validates the data collection and post-processing chain used in all later experiments.
- **Measure shot noise as a function of LO power (A2).** Step the laser drive current to vary the LO power reaching the balanced detector, and confirm that the excess noise above the electronic floor grows linearly with LO power, as expected for shot noise.
- **Choose the operating point.** From the A2 results, find the LO power at which shot noise overtakes the electronic floor, and choose a laser operating point at which the system is close enough to the shot-noise limit to meet requirement P2 (Table 3.1).
- **Quantify RIN and validate balanced detection (A3).** Repeat the measurement with only one photodiode connected, making the PD100B a single-ended detector. Comparing this with the balanced case measures the laser's relative intensity noise (RIN) and confirms that the balanced detector cancels this common-mode noise as intended.

4.2.2 Experimental Methodology

The experimental setup used is given in Figure 4.1:

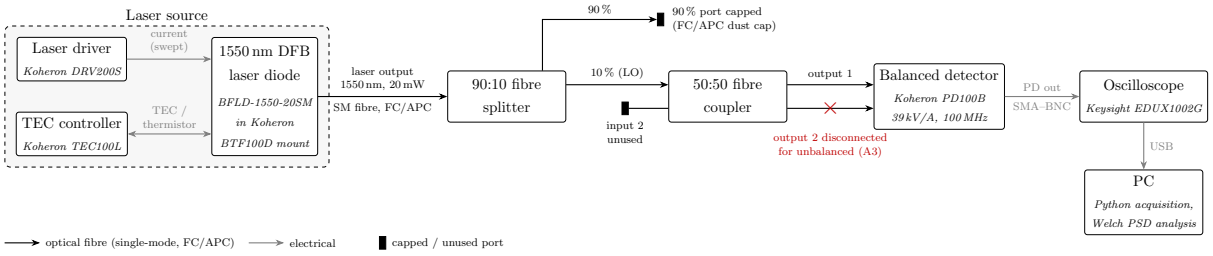


Figure 4.1: Experimental setup A

It can be seen that this differs from the full system setup. The circulator and collimator are unused for this experiment. The 90% splitter arm is left unconnected and is capped. Only the 10% splitter arm (LO) is connected to the one arm of the 50:50 coupler. Both outputs of the coupler are connected to the balanced detector, except in the case of

experiment A3. Since there is no circulator, collimator or target, the only light reaching the detector is the LO itself. With no second optical field there is no beat term (as in Equation 2.4), and the LO contributes only a DC photocurrent, which is removed by the detector's AC coupling. Thus, the signal at the detector will be a pure noise signal, containing only electronic noise and shot noise. Note that phase-induced intensity noise is not present here, as it originates from reflections in the collimator and circulator which are not present in this setup, and RIN is assumed to be cancelled fully by balanced detection (further investigated in A3).

The experiment is divided into three phases, with slightly different configurations for each in order to give different measurands, as shown in Table 4.1:

ID	**Configuration**	**Quantity Measured**
A1	Laser output off, detector switched on	Electronic noise only
A2	Laser output on (current varied), detector switched on, both detector inputs connected	Electronic noise + shot noise
A3	Laser output on (current varied), detector switched on, only one detector input connected	Electronic noise + shot noise + full RIN

Table 4.1: Experiment A phases and configurations

For each measurement, the laser is held at a fixed drive current and 100 records of the detector output are captured sequentially with the oscilloscope and transferred to a laptop. Each record contains 25 000 samples at 125 MSa/s (200 μs). A1 requires only a single measurement, while A2 and A3 are repeated over a range of drive currents up to the 48 mA operating point, with A3 performed on each detector input in turn.

The PSD of each record is estimated using Welch's method (Hann window, 2 500-sample segments with 50 % overlap, giving a 50 kHz resolution) rather than a single periodogram, since averaging over segments greatly reduces the variance of the estimate, giving a smoother curve. Because the noise is stationary (properties don't change across records), the PSDs of the 100 records can then be averaged as well for further smoothing.

To compare the results with a single figure, the PSD can be averaged over the 5–20 MHz band. From Equation 2.1, this band corresponds to target speeds of roughly 4–15.5 m/s, making it a reasonable band to take an average noise power measurement over, as it is where the Doppler peaks in subsequent tests would appear.

4.2.3 Results and Discussion

A1: Electronic Noise

The power spectral density plot in Figure 4.2 was obtained from measurement A1, which gives the electronic noise from the detector only:

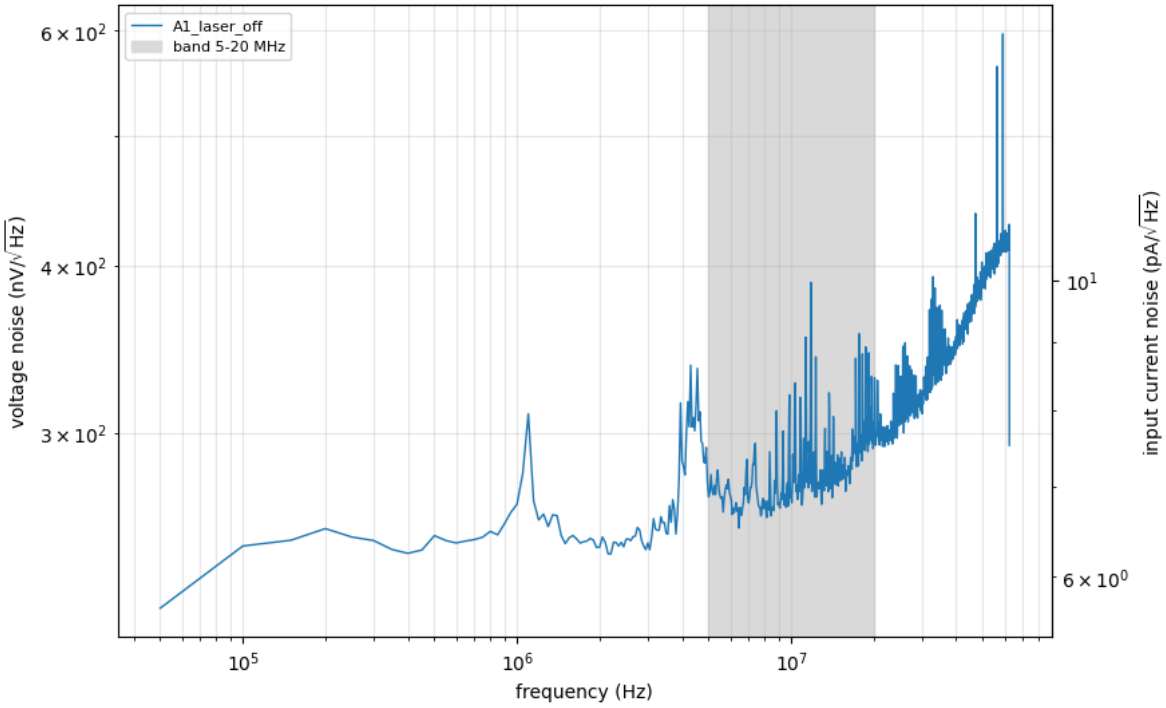


Figure 4.2: A1: PSD of electronic noise from detector

The obtained PSD plot is noisy, despite the use of the Welch PSD and averaging, however it is sufficient for the purposes of this investigation. It can be seen that the electronic noise increases with frequency. The average PSD over the 5 - 20MHz band is calculated (using the script) as: $8.06 \times 10^{-14} V^2/Hz$. Using the PD100B detector's gain value of 39kV/A, this can be converted into an average input noise current density across 5 - 20MHz, for comparison to the PD100B datasheet:

$$I_{input} = \frac{8.06 \times 10^{-14} V^2/Hz}{(39kV/A)^2} = 7.28pA/\sqrt{Hz} \tag{4.1}$$

Input noise current density can be seen as the equivalent of injecting this noise current

at the detector input. The calculated average of $7.28pA/\sqrt{Hz}$ across 5 - 20MHz agrees closely with the PD100B's datasheet value of $8pA/\sqrt{Hz}$, measured at 10MHz [**?**]. Thus, this level of electronic noise is expected, and the measurement and post processing chain is validated.

In order to achieve shot-noise limited operation, the average PSD of the shot-noise over this region must thus be $\gg 8.06 \times 10^{-14} V^2/Hz$.

A2: Electronic Noise & Shot Noise

Figure 4.3 gives the PSD plots of the total noise - made up of shot-noise, electronic noise, and residual RIN - for different laser currents on the same axes:

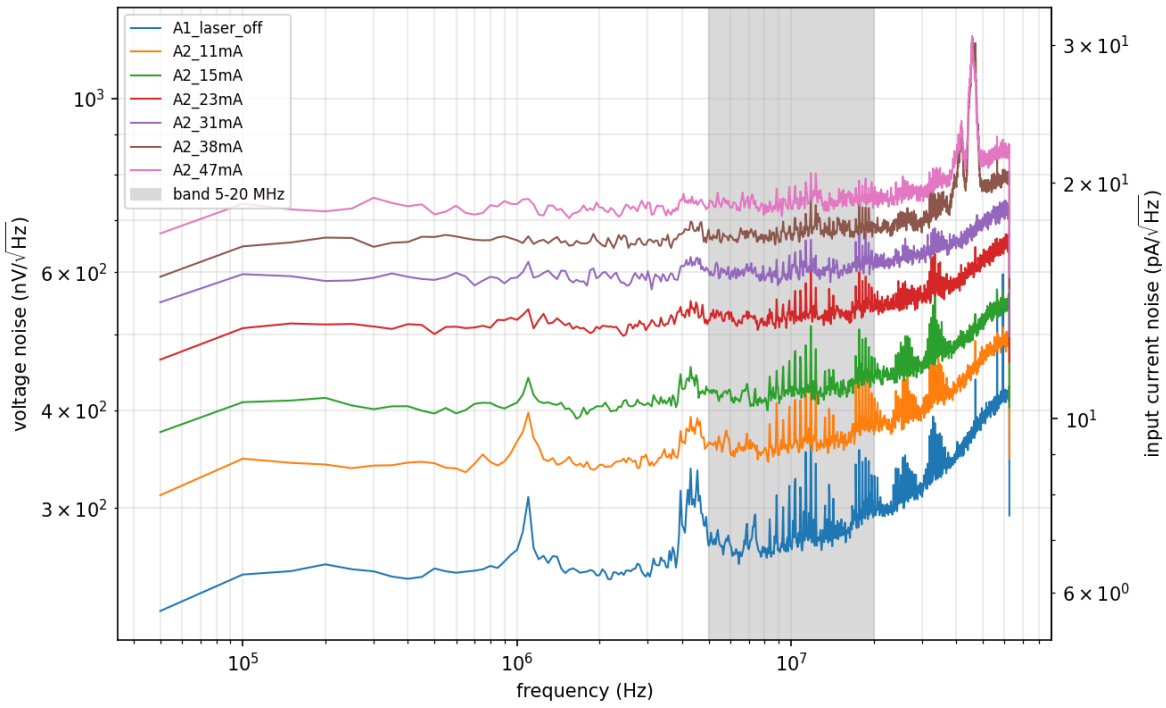


Figure 4.3: PSD of total noise (shot, electronic and relative intensity noise) for different laser source currents

It can be seen that the total noise increases with increased laser current - this is expected, due to the linear relationship between shot noise and LO power (which is related to laser current which is varied here). In order to compare the PSDs for the various laser currents, the average PSD across the 5 - 20MHz range can again be taken here, for each laser current. Figure 4.4 gives these averages plotted against laser current:

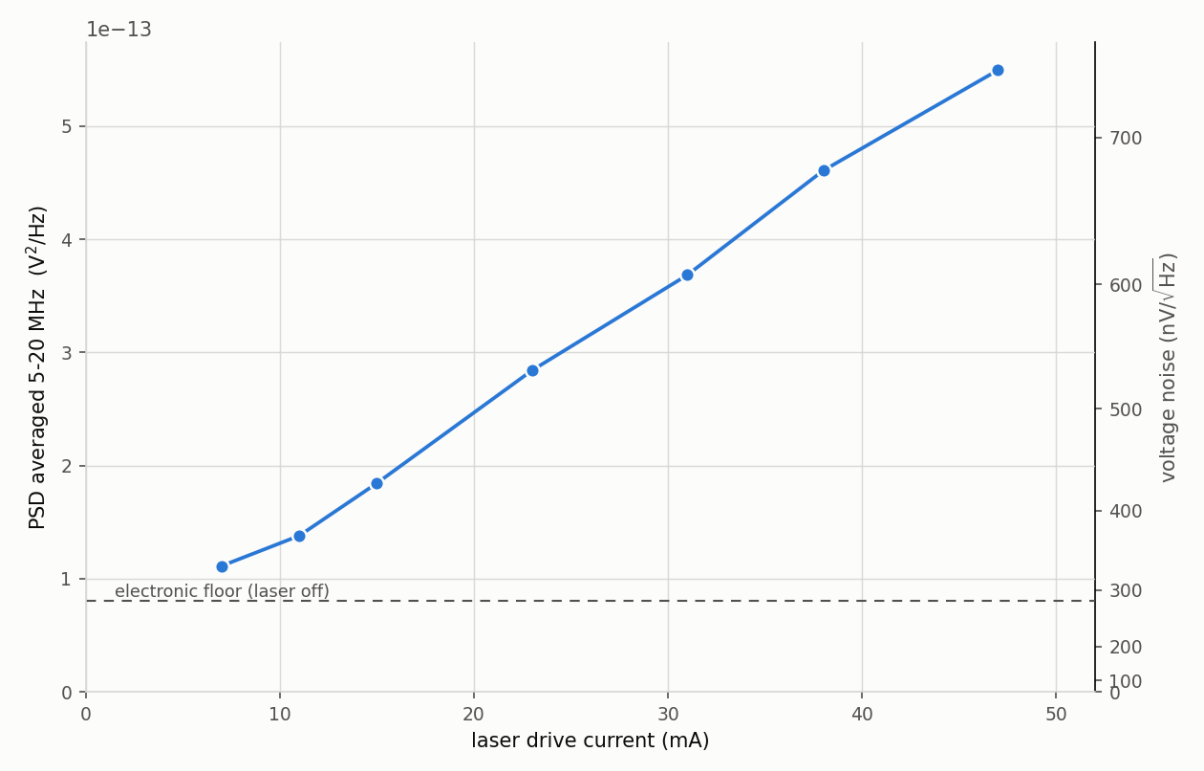


Figure 4.4: 5 - 20MHz Average PSD vs Laser Current with Balanced Detection

It can clearly be seen in Figure 4.4 and Figure 4.3 that the shot-noise exceeds the electronic noise floor at even the lowest measured laser currents, thus making the system shot-noise limited. As the current is increased, the system approaches the ideal shot-noise limit, as shot noise increasingly swamps out the electronic noise floor.

Performance requirement P2 in Table 3.1 requires that the total detector noise is within 1dB of the shot-noise limit. This requirement can be decomposed as follows:

Noise power adds in linear units, so letting S be shot noise, and E be every other noise source present here (electronic noise and any remaining RIN), the total noise floor, T is given by:

$$T = S + E = S(1 + \frac{E}{S}) \tag{4.2}$$

Converting to dB:

$$10log_{10}(S(1 + \frac{E}{S})) = 10log_{10}(S) + 10log_{10}(1 + \frac{E}{S}) \tag{4.3}$$

In this form, it can clearly be seen that as S increases, the effect of E on the total noise floor is increasingly reduced, as the second term, $10log_{10}(1 + \frac{E}{S})$ approaches zero. This second term is the excess noise above the ideal shot-noise limit, which is given by $10log_{10}(S)$. Thus, in order to obtain a total noise floor which is within 1dB of the shot noise limit, we set:

$$10log_{10}(1 + \frac{E}{S}) = 1dB \tag{4.4}$$

$$\rightarrow \frac{E}{S} \leq 0.26 \tag{4.5}$$

Substituting the obtained electronic noise floor obtained in A1 into Equation 4.5, we obtain:

$$S \geq 3.1 \times 10^{-13} V^2/Hz \tag{4.6}$$

Noting that Figure 4.3 and Figure 4.4 both give the **total** detector noise floor (shot **and** electronic noise), we require the total noise to be above $3.906 \times 10^{-13} V^2/Hz$ to gives a shot noise value of $3.1 \times 10^{-13} V^2/Hz$. This can be seen to be reached at a laser current of approximately 32mA. Above 32mA, the total detector noise floor is within 1dB of the ideal shot-noise limit, and the ideal shot-noise limit is further approached as the current is increased.

From this, it can be seen that an operating current of 48mA for the laser allows P2 in Table 3.1 to be met with the highest margin, and this will thus be used as the laser current for the remainder of this project. While a higher laser current still may be used, to even further approach the ideal shot-noise limit, this risks saturating the detector. The detector has a maximum incident power rating of 1.5mW on each photodiode [**?**]. From the laser datasheet [**?**], a current of 48mA gives approximately 17mW at the laser output, which will result in approximately 0.85mW at each detector from the LO alone. This allows the detector to thus remain comfortably far away from its saturation point, even if returned signal is high and results in a higher incident power at the photodiode after mixing with the LO.

A3: Relative Intensity Noise

Experiment A3 serves to illustrate the effectiveness of balanced detection in removing relative intensity noise at the detector output. Figure 4.5 gives the 5 - 20MHz averaged PSD of the detected noise, when only a single photodetector input of the PD100B is used - effectively turning it into a singe-ended detector. The electronic noise floor will remain the same, but the total noise PSD averages will now also contain RIN in addition to shot-noise and electronic noise. Before analysis of this plot can commence, it is important to note that in the case of balanced detection, the shot-noise of each detector is uncorrelated, and will therefore add together at the output. Thus, in the single-ended case, the shot-noise at the output for a given laser current is expected to be half that of the balanced detection case. Any excess noise above this half-power shot-noise level is thus contributed from the electronic noise, and the uncancelled RIN.

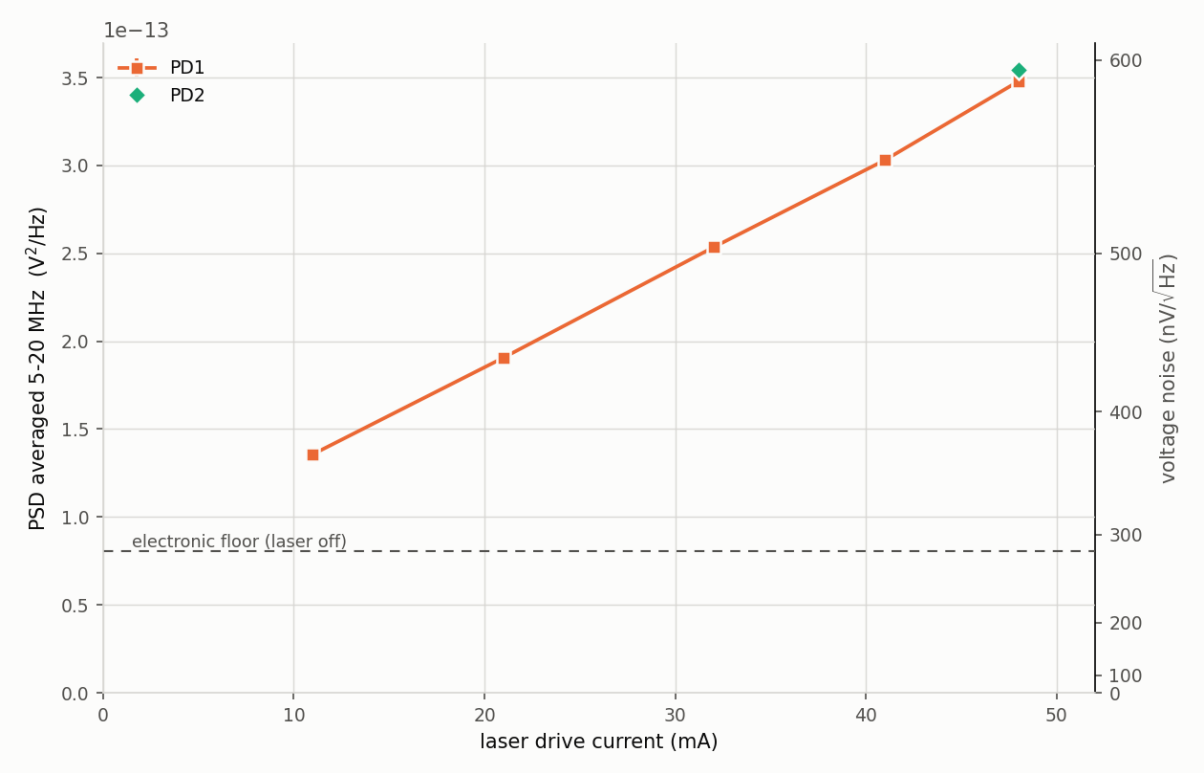


Figure 4.5: A3: 5 - 20MHz Average PSD vs Laser Current with Single-Ended Detection

It can be seen in Figure 4.5 that indeed, the total noise is lower than in Figure 4.4, as expected. Observing the total noise at the operating point of 48mA however, it can be seen that a value of approximately $3.5 \times 10^{-13} V^2/Hz$ is obtained, which gives $2.694 \times 10^{-13} V^2/Hz$ once electronic noise has been subtracted. In the balanced detection

case at the same laser current, a value of $5.49 \times 10^{-13} V^2/Hz$ was obtained - subtracting electronic noise and halving this gives $2.342 \times 10^{-13} V^2/Hz$. This clearly indicates that the noise power in the single-ended case is higher than simply half of the shot-noise in the balanced detection case, thus proving that indeed there is excess RIN present. Thisexcess RIN noise can be calculated:

$$RIN_{excess} = (2.694 - 2.342) \times 10^{-13} = 0.352 \times 10^{-13} V^2/Hz \tag{4.7}$$

Converting to dB:

$$10log_{10}(1 + \frac{RIN_{excess}}{S}) = 10log_{10}(1 + \frac{0.352 \times 10^{-13}}{2.324 \times 10^{-13}}) = 0.6125dB \tag{4.8}$$

Thus, RIN contributes an additional 0.6125dB of noise above the ideal shot-noise limit at the 48mA operating point. Balanced detection allows this excess noise to be removed, as illustrated in the results of A2 - allowing for a better SNR.

4.3 B: Phase-induced Intensity Noise

4.3.1 Experimental Goals

As discussed in subsection 2.3.5, and elaborated in section A.1, Phase-induced Intensity Noise (PIIN) can be a significant contributor to the total noise of the system, rising far above shot-noise - particularly when the linewidth of the laser source is not narrow, and when there is significant path-length mismatch between the LO and the transmit signal.

The goal of Experiment B is to measure PIIN in a controlled setting, with no target, thus making all PIIN in the measurement attributable to reflections from the circulator and collimator, and independent of target.

4.3.2 Experimental Methodology

Chapter 5

Results

These are the results I found from my investigation.

Present your results in a suitable format using tables and graphs where necessary. Remember to refer to them in text and caption them properly.

5.1 Simulation Results

5.2 Experimental Results

Chapter 6

Discussion

Here is what the results mean and how they tie to existing literature...

Discuss the relevance of your results and how they fit into the theoretical work you described in your literature review.

Chapter 7

Conclusions

These are the conclusions from the investivation and how the investigation changes things in this field or contributes to current knowledge...

Draw suitable and intelligent conclusions from your results and subsequent discussion.

Chapter 8

Recommendations

Make sensible recommendations for further work.

Bibliography

[1] Cyrus F. Abari, Anders T. Pedersen, and Jakob Mann. An all-fiber image-reject homodyne coherent Doppler wind lidar. *Optics Express*, 22(21):25880–25894, October 2014.

[2] Sameh Abdelazim, David Santoro, Mark F. Arend, Fred Moshary, and Sam Ahmed. Development and Operational Analysis of an All-Fiber Coherent Doppler Lidar System for Wind Sensing and Aerosol Profiling. *IEEE Transactions on Geoscience and Remote Sensing*, 53(12):6495–6506, December 2015.

[3] Michael Bennett. Doppler lidar for boundary-layer measurements: must it be expensive? page 249, Varna, Bulgaria, November 2003.

[4] Tony L. Burton, Nick Jenkins, Ervin Bossanyi, David Sharpe, and Michael Graham. *Wind Energy Handbook*. John Wiley & Sons, Chichester, UK, 3rd edition, 2021.

[5] Jean-Pierre Cariou. LEOSPHERE Pulsed Lidar Principles.

[6] Jean-Pierre Cariou, Béatrice Augere, and Matthieu Valla. Laser source requirements for coherent lidars based on fiber technology. *Comptes Rendus Physique*, 7(2):213–223, March 2006. High power fiber lasers and amplifiers.

[7] Luke Colberg, Dylan J. Maxwell, and Kevin S. Repasky. Diode-laser-based direct-detection Doppler wind lidar: design and initial results. *Applied Optics*, 65(10):3284–3299, April 2026.

[8] Christian Doppler. Über das farbige Licht der Doppelsterne und einiger anderer Gestirne des Himmels. *Abhandlungen der Königlichen Böhmischen Gesellschaft der Wissenschaften*, 2:465–482, 1842. Read before the Royal Bohemian Society of Sciences, Prague, 25 May 1842.

[9] Global Wind Energy Council. Global wind report 2026. Technical report, Global Wind Energy Council, April 2026. Accessed DD Month 2026.

[10] M. Harris, G. N. Pearson, J. M. Vaughan, D. Letalick, and C. Karlsson. The role of laser coherence length in continuous-wave coherent laser radar. *Journal of Modern Optics*, 45(8):1567–1581, August 1998.

[11] Michael Harris, Maureen M. Hand, and Alan D. Wright. Lidar for turbine control. Technical Report NREL/TP-500-39154, National Renewable Energy Laboratory, Golden, CO, January 2006.

[12] Chris Hill. Remote Sensing (UpWind WP6) QinetiQ Lidar Measurement Report.

[13] International Electrotechnical Commission. Wind energy generation systems – part 50-3: Use of nacelle-mounted lidars for wind measurements. International Standard IEC 61400-50-3:2022, International Electrotechnical Commission, Geneva, January 2022. Edition 1.0.

[14] Shan Jiang, Dongsong Sun, Yuli Han, Fei Han, Anran Zhou, and Jun Zheng. Performance of continuous-wave coherent doppler lidar for wind measurement. *Current Optics and Photonics*, 3(5):466–472, oct 2019.

[15] Kazuro Kikuchi. Fundamentals of coherent optical fiber communications. *Journal of Lightwave Technology*, 34(1):157–179, jan 2016. Tutorial Review.

[16] C. Laurence Korb, Bruce M. Gentry, S. Xingfu Li, and Cristina Flesia. Theory of the double-edge technique for Doppler lidar wind measurement. *Applied Optics*, 37(15):3097–3104, May 1998.

[17] Stefano Letizia, Peter Brugger, Nicola Bodini, Raghavendra Krishnamurthy, Andrew Scholbrock, Eric Simley, Fernando Porté-Agel, Nicholas Hamilton, Paula Doubrawa, and Patrick Moriarty. Characterization of wind turbine flow through nacelle-mounted lidars: a review. *Frontiers in Mechanical Engineering*, 9, November 2023.

[18] Zhengliang Liu, Janet F. Barlow, Pak-Wai Chan, Jimmy Chi Hung Fung, Yuguo Li, Chao Ren, Hugo Wai Leung Mak, and Edward Ng. A Review of Progress and Applications of Pulsed Doppler Wind LiDARs. *Remote Sensing*, 11(21):2522, January 2019.

[19] T. Mikkelsen, N. Angelou, K. Hansen, M. Sjöholm, M. Harris, C. Slinger, P. Hadley, R. Scullion, G. Ellis, and G. Vives. A spinner-integrated wind lidar for enhanced wind turbine control. *Wind Energy*, 16(4):625–643, 2013. _eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/we.1564.

[20] Rüdiger Paschotta. Shot noise. RP Photonics Encyclopedia, 2005. Accessed: 2026-09-18.

[21] Eric Simley, Holger Fürst, Florian Haizmann, and David Schlipf. Optimizing Lidars for Wind Turbine Control Applications—Results from the IEA Wind Task 32 Workshop. *Remote Sensing*, 10(6):863, June 2018.

[22] Paul Veers, Katherine Dykes, Eric Lantz, Stephan Barth, Carlo L. Bottasso, Ola Carlson, Andrew Clifton, Johney Green, Peter Green, Hannele Holttinen, Daniel Laird, Ville Lehtomäki, Julie K. Lundquist, James Manwell, Melinda Marquis, Charles Meneveau, Patrick Moriarty, Xabier Munduate, Michael Muskulus, Jonathan Naughton, Lucy Pao, Joshua Paquette, Joachim Peinke, Amy Robertson, Javier Sanz Rodrigo, Anna Maria Sempreviva, J. Charles Smith, Aidan Tuohy, and Ryan Wiser. Grand challenges in the science of wind energy. *Science*, 366(6464):eaau2027, October 2019.

[23] R. Wagner, T.f. Pedersen, M. Courtney, I. Antoniou, S. Davoust, and R.l. Rivera. Power curve measurement with a nacelle mounted lidar. *Wind Energy*, 17(9):1441–1453, 2014. _eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/we.1643.

[24] Rozenn Wagner, Michael Stephen Courtney, Troels F. Pedersen, and Samuel Davoust. Uncertainty of power curve measurement with a two-beam nacelle-mounted lidar. *Wind Energy*, 19(7):1269–1287, 2016. _eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1002/we.1897.

[25] Christian Werner. Doppler wind lidar. In Claus Weitkamp, editor, *Lidar: Range-Resolved Optical Remote Sensing of the Atmosphere*, volume 102 of *Springer Series in Optical Sciences*, pages 325–354. Springer, New York, NY, 2005.

Appendix A

Appendix A

A.1 Linewidth, PIIN and Doppler peak broadening

In order to understand the linewidth property of a laser source, it is important to first have an understanding of the concept of phase noise in a laser:

An ideal laser emits an electromagnetic wave $E(t) = E_0 cos(2\pi f_0 t + \phi)$ with ϕ a constant. A real laser emits $E(t) = E_0 cos(2\pi f_0 t + \phi(t))$ where $\phi(t)$ drifts unpredictably. This drift is called phase noise. Thus, when considering the frequency spectrum of the emitted laser beam over a long time period, it takes a Lorentzian shape centred around the nominal frequency of the laser (whereas an ideal laser would present simply as a delta spike at the nominal frequency). The laser linewidth, Δf, is then defined as the -3dB bandwidth of this Lorentzian spectrum, as illustrated in Figure A.1 below:

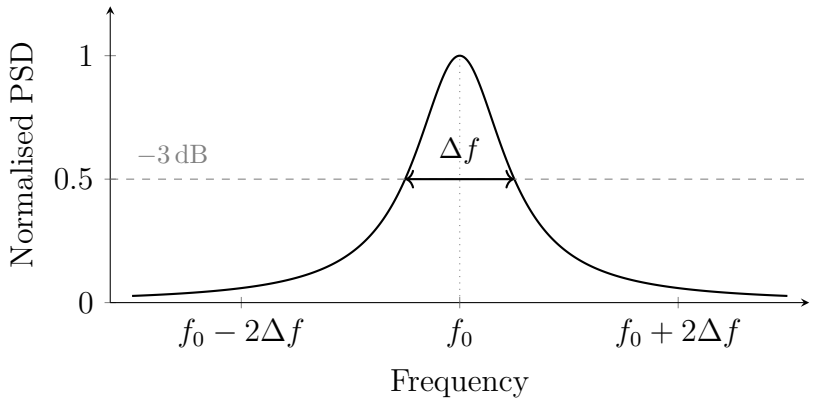


Figure A.1: Lorentzian frequency spectrum of a laser diode.

Asssociated with the linewidth is a quantity termed coherence length. The coherence length of the emitted beam, L_c, is given by [3]:

$$L_c = \frac{c}{2\pi\Delta f} \tag{A.1}$$

Where Δf is the linewidth of the source laser, and c is the speed of light. The coherence length is of relevance when considering linewidth effects.

The non-ideal linewidth property of a laser source gives rise to two primary effects in a DWL system's measurements - namely phase-induced intensity noise (PIIN), and the broadened Doppler peak:

As noted by Cariou et al. [6], Jiang et al. [14], and Harris et al. [10], reflections of the outgoing beam - caused by optics and the fibre optic transmit path - travel back in the system and are interfered with the reference beam (LO) before detection (in the same way as the actual backscattered signal of interest is). Because the reflection and LO derive from the same source, the mixing product is a strong, unwanted tone at baseband (homodyne) or at the intermediate frequency (heterodyne), carrying the phase noise accumulated over the path-length mismatch ΔL between the reflection and the LO. Harris et al. [10] show that the form of this residual noise is set by how large ΔL is relative to the coherence length (L_c) of the laser source.

For $\Delta L \gg L_c$ the two phases are highly uncorrelated because the difference in path (and therefore time) taken has allowed the phase of the two to "wander" significantly - and thus the tone broadens into a Lorentzian of width $2\Delta f$. This case is illustrated in Figure A.2 below, where it can be seen that this unwanted Lorentzian peak can serve to swamp out the desired Doppler peak to be measured - this is especially impactful for low wind-speeds, where the Doppler peak will be closer to the unwanted common-mode Lorentzian peak, and will be hidden more by it.

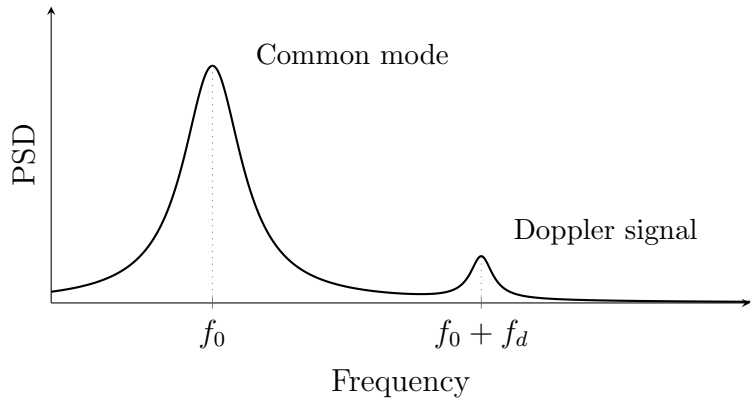


Figure A.2: Spectrum of a monostatic CW coherent lidar, showing the Doppler signal superimposed on the skirt of the common-mode peak (adapted from [6]).

For $\Delta L \ll L_c$, the case where the mismatch between LO and reflection paths is smaller, the reflected signal and the LO will be more correlated in phase. Thus, the residual instead appears as a flat excess noise floor across the frequency band, of level proportional to $\Delta f\, \Delta L^2$ [10]. This excess noise is termed excess LO noise, interferometric noise or phase-induced intensity noise (PIIN), depending on author [10, 1, 14]. For the remainder of this paper, this excess noise will be termed PIIN. PIIN has been measured at 5 - 30dB above shot noise in a diode-laser lidar [10] and well above shot noise in a fibre-laser system [14]. Thus, PIIN is a significant noise source which can cause the system to operate far from the ideal shot-noise limit discussed in subsection 2.3.5. It is thus clear that to reduce PIIN, a laser with as small a linewidth as possible should be selected, and the path length mismatch, $\Delta L\ L$ between LO and reflection paths should be minimised.

Another consideration arising due to non-ideal laser linewidth is the broadening of the measurable Doppler peak. This is particularly notable for longer ranges.

If the coherence length from Equation A.1 is long relative to the emitted signal's path length, the broadening of the Doppler peak due to the laser's linewidth may be neglected, and it will instead be other factors (such as turbulence in the air, and speckle) which dominate and cause a broader measurable peak [3]. This is because the total round-trip time of the beam is shorter, thus ensuring that the phase of the beam has not had "time" to shift significantly in between transmission, and later reception and mixing with the local oscillator. Thus, the received backscatter signal and the LO phase will be more coherent and will result in a sharper peak after mixing.

Conversely, if the coherence length is short relative to the signal path length (i.e. for targets at greater ranges), the Doppler peak will be spread over $2 \times \Delta f$, meaning that

the measurement resolution and the ability to resolve turbulence in the air will be reduced [3].

From the two effects described above, it can clearly be seen that a narrow linewidth laser source is highly beneficial for DWL systems, in order to achieve the best-case shot-noise limited operation, and to achieve a narrow, measurable Doppler peak.

IMPORTANT: Appendix A (see the table below) should provide a summary of how you have met the GAs associated with the course, and where in the report the evidence can be found.

If you have used AI in writing your report, you need to provide details in one of the appendices of how you used it.

If appropriate, in a subsequent appendix, provide a link to any GitHub repository where the code and additional materials for your project can be found.

GA	**Requirement**	**Justification and section in the report**
1	Problem-solving	Description here
4	Investigations, experiments and data analysis	Description here
5	Use of engineering tools	Description here
6	Professional and technical communication (Long report)	Description here
8	Individual work	Description here
9	Independent learning ability	Description here

Appendix B

Addenda

B.1 Ethics Forms